

The Linear Model (LM)

T. Florian Jaeger

September 6, 2024

Readings and assignments in preparation of this class

In preparation for class, please **read James et al. (2013), Chapter 3 up to—but not including—3.2 (12 pages)**. This text is quite intuitive and assumes comparatively little background. But if you find yourself getting lost, please stop and first read or watch one of the many more informal introductions to linear models (e.g., this one). Then try to read James et al. (2013).¹ **You should plan for 1-2h of reading and summarizing your questions in preparation for this week of the class.**

Optionally, those of you have interests beyond the basics of the LM, and find it intriguing to ponder about why we use the models we use, might enjoy reading p. 73-78 of McElreath's (2019) *Statistical Rethinking*². It's an amazing book, and these few pages accessibly and thought-provokingly take you through the question of why we often feel comfortable assuming normality in our analyses (as is done in the linear model).

Also optionally, you might find it helpful to read and *work through* this document. **Please note that this document is providing R code only.** Some of the steps described here might have less direct solutions in Matlab or Python, so it is recommended that you start early.

¹ James et al. (2013) *An introduction to statistical learning* is the an attempt—a wonderful attempt, full of effective visualization, I think—to make the content of one of the most influential books in machine learning (“Statistical learning” by Tibshirani and colleagues, the developers of bootstrap and other important non-parametric approaches) accessible to a non-expert audience.

About this document

This PDF is generated from an R markdown document (using the `tufte` package for styling). If you're interested in understanding the code that generated the examples presented below, you can clone the repo. Both the PDF and the R markdown that the PDF is generated from are in the `scripts/` folder. You can go through the document line by line, or make changes and ‘knit’ a new PDF with the press of a button.

The linear model (LM)

LMS (aka linear regression) can be thought of as an example of a much larger family of models that describe an *outcome* variable y (aka the *dependent* variable) as the function g of one or more *predictor* variables x_1, \dots, x_k (aka *independent* variables), $y = g(x_1, \dots, x_k)$. In this sense, the LM is a—very constrained—*predictive model*. Specifi-

cally, the LM assumes that y is a draw from a Normal distribution \mathcal{N} with constant standard deviation σ and a mean μ that is a *linear combination* (the weighted sum) of all the predictors included in the model:

$$\mathbf{y} \sim \mathcal{N}(\mu, \sigma_{resid}) \quad (1)$$

$$\mu = \beta_1 \mathbf{x}_1 + \dots \beta_k \mathbf{x}_k \quad (2)$$

$$= X\beta \quad (3)$$

Here, y is a column vector of outcome observations of length n (e.g., a column in your data table), X is called the $n \times k$ *model (or design) matrix* (e.g., in the simplest case, a number of k columns in your data table) and β is a vector of weight values of length k (aka coefficients)—one for each of the k predictors in the model.

In the literature, LMs are often written in a number of different ways. For instance, one might describe the linear model for each observation i of the outcome y_i and all predictors $x_{1,i}, \dots, x_{k,i}$:²

$$y_i = \beta_1 x_{1,i} + \dots + \beta_k x_{k,i} + \epsilon_i, \epsilon_i \sim \text{Normal}(0, \sigma_{resid}) \quad (4)$$

or in vector form, or for as in James et al. (2013, p. 63):

$$\mathbf{y} = \beta_1 \mathbf{x}_1 + \dots + \beta_k \mathbf{x}_k + \epsilon, \epsilon \sim \text{Normal}(0, \sigma_{resid}) \quad (5)$$

$$= X\beta + \epsilon, \epsilon \sim \text{Normal}(0, \sigma_{resid}) \quad (6)$$

or combining equations (1) and (3) from above, as in Gelman & Hill (2007, p. 38):

$$\mathbf{y} = \text{Normal}(X\beta, \sigma_{resid}) \quad (7)$$

Any of these formulations are equivalent, and all are used in the literature. Finally, a common—but arbitrary—convention is to use a different symbol for one particular predictor, the *intercept* α :

$$y \sim \mathcal{N}(\mu, \sigma) \quad (8)$$

$$\mu = \alpha + \beta_1 x_1 + \dots \beta_k x_k \quad (9)$$

$$= \alpha x_0 + \beta_1 x_1 + \dots \beta_k x_k \quad (10)$$

Here, x_0 is assumed to be the constant 1, i.e., $x_0 = 1$. This convention singles out α , which helps to realize that it has an intuitive interpretation: it is the predicted value of the outcome y when all $x_1 = \dots = x_k = 0$.




Figure 1: A graphical model visualization of the linear model. Shaded circles are observable variables. White circles are latent variables. These are the variables that constitute the degrees of freedom of the model, the parameters that are fit to the data. Diamonds shapes indicate variables that are determined by other variables. The index $i = 1, \dots, N$ refers identifies the observations (usually, the rows) in your data. Note that σ , unlike μ_i is assumed to be constant across observations. This is the assumption of the homogeneity of variance.

² ϵ_i is the *residual* for observation i (observed y_i - predicted \hat{y}_i).

Assumptions of the LM

As a theoretical model the LM makes several assumptions (cf. Gelman & Hill, 2007, p. 45). The first two assumptions are shared with many/most other statistical models we use (to be precise, some models don't strictly make these assumptions, but they do correct for their violations):

- **Independence of observations/errors.** Each observation y is assumed to be drawn independent of all other observations.
- **Validity/exhaustivity.**

The other assumptions are more specific to the LM, though the last two of these are shared with GLMs and GLMMs, as we'll see in later classes:³

- **Normality of errors.** y is drawn from a Normal distribution \mathcal{N} (with mean $\mu = X\beta$), and thus is assumed to be able to take any real value from $-\infty$ to $+\infty$. Equivalently, we can say that the *residuals* (or prediction errors) are drawn from a Normal distribution (with mean 0).
- **Equal variance of errors.** While $E[y] = \bar{y}$ is assumed to vary with the predictors X , the residual standard deviation around that mean σ is assumed to be independent of X , and thus constant and identical across all observations.
- **Additivity of effects.** The *effects* of predictors are assumed to be additive: $E[y]$ is assumed to be the weighted *sum* of the predictors.
- **Linearity of each effects.** y is linear in the β s. That is, the *effects* of predictors are assumed to be linear.

Many, if not most, of these assumptions are actually violated when we apply the model. To some extent linear regression is robust to such violations, depending on the specific assumption. For example, instead of equality of variance it is often sufficient that variances are similar (*homogeneity* of variance) or that the variance-covariance matrix does not exhibit *heteroscedasticity*.⁴ **But violations of the independence, normality, or homoscedasticity assumptions can make a model invalid, and the statistical inferences based on it invalid** (e.g., by inflating Type I and II errors). We return to these points in later classes.

Use of the LM in our field(s)

While the LM can be used as a predictive model—i.e., to make predictions about unseen data—this is rarely done in our fields. Rather, we tend to fit an LM to our data, and then conduct statistical inferences about the *effects* of one or more of the model's predictors based

³ The last two assumptions often confuse learners. To understand them it is important to distinguish between **input variables** and the **predictors** we derive from them. For example, in the case study presented below, one of the input variables is *VOT*. But we could include both *VOT* and *VOT*² as predictors in the model. This would model both linear and quadratic effects of *VOT*, each of which is assumed linearly add to the predicted outcome y . Similarly, we can include interactions of the input variable *VOT* and some other predictor in the model, which will add their product as a predictor to the model. That is, we can very much use an LM to ask whether two variables have super- or sub-additive interactions. The additivity assumption just means that the effects of the interaction predictor are assumed to be added to the effects of the other predictors in the model.

⁴ From Ancient Greek *hetero*—"different" and *skedasis*—"dispersion". Homoscedasticity extends the notion of homogeneity of variance to the covariance between variables

on the best-fitting parameterization of that model. Best-fitting parameterization refers to the choice of β s (and the residual standard deviation σ , though we don't tend to be interested in this parameter) that make the LM deliver the best predictions about the observed values of y .

We'll learn later what "best" means here, and how those best-fitting parameters and their standard errors are obtained. For now, we'll assume that those are known, and focus on what they tell us about the effects of the different predictors that are included in the LM.

When we employ an LM during data analysis, we are making the assumption that this model describes the relation between the predictors and the outcome in the population we seek to study. This also means that all of our conclusions are contingent on these assumptions, including in particular the assumption that we have considered all relevant variables (or have otherwise ruled out that other effects can confound our analysis).

Geometric intuitions

I won't repeat the geometric interpretation of LMs here, as the James et al. readings does a great job setting that up. Generally speaking, an LM with k predictors (plus the intercept) can be thought of as defining a hyperplane in a $k + 1$ dimensional space, where the additional dimension is the outcome. For instance, for a single predictor, the LM defines a line that describes the predicted outcome as a function of the single predictor. For two predictors, the LM defines a plane in the 3D space defined by the two predictors and the outcome, etc.

Often it is more intuitive, however, to plot the data in fewer dimensions. For example, an LM with a continuous predictor, a categorical predictor (factor), and their interaction, describes multiple lines (that might differ in slope). An LM with a quadratic polynomial of a single continuous input variable will have two continuous predictors derived from that input variable. But rather than to plot the predictions in 3D space, we can plot them in 2D space, showing the effects of both the linear and the quadratic transforms of the input variable along a single axis.

General workflow

Here's the general workflow I'd recommend for an LM analysis, following adequate data checks, variable typing (e.g., making sure that variables that are factors indeed are factors, etc.) and filtering

(e.g., removal of practice and/or catch trials, outlier exclusion, etc.):

- Get data into table format, with each observation of your data corresponding to one row of the table and separate columns for the outcome and each of the input variables.⁵
- Decide how to code the factors in your data, so that each of the factor input variables can be translated into one or more numerical predictors.⁶ Decide whether to transform your continuous input variables, e.g., by centering and/or standardizing them.
- Define the regression formula $y \sim 1 + x_1 + \dots + x_k$, where 1 stands for the constant 1. Together with the data table and information about the factor coding (which is stored as an attribute to the factor column in the data.frame in R and in many similar programs), this formula contains all the information that is required for a function intended to fit the model to the data.

A call to the LM-fitting function typically first implicitly creates the model matrix X . **Looking at the model matrix corresponding to an LM or other model can be a very effective way to learn to understand formula syntax in whatever language you're using.**

By default, this model matrix will contain the constant 1 as the first column, followed by one column per *predictor*. Note that this can be more columns than the number of input variables. For example, factor input variables with k distinct levels will result in $k - 1$ distinct predictors, and thus $k - 1$ distinct columns in the model matrix. Similarly, interactions in the formula add additional columns to the model matrix (more on that in a later lecture). And there are a number of other common formula abbreviations that will expand into multiple columns in the design matrix. If you're not sure what a formula does, I recommend simply looking at the model matrix it creates. E.g., in R:

```
model.matrix(lm(data = d, formula = Response.RT ~ 1 + Condition * VOT)) %>%
  # Select 20 random rows from the model matrix for illustration
  as_tibble() %>%
  slice_sample(n = 15)
```

```
# A tibble: 15 x 6
  '(Intercept)' ConditionShift10 ConditionShift40 VOT 'ConditionShift10:VOT' 'ConditionShift40:VOT'
    <dbl>          <dbl>          <dbl> <dbl>          <dbl>          <dbl>
1         1             0             0    15             0             0
2         1             0             0     5             0             0
3         1             0             0     5             0             0
4         1             0             1    35             0            35
5         1             0             0    55             0             0
6         1             0             1    40             0            40
7         1             0             1    75             0            75
8         1             1             0     5             5             0
9         1             0             0    75             0             0
```

⁵ What counts as an observation will depend on both your hypothesis (e.g., whether it is meant to test a hypothesis within a participant or across participants) and the type of model you're using (e.g., an LM vs. a linear mixed-effects model).

⁶ In R, the `contrasts()` function and the family of `contr.*` functions are very helpful in setting up the contrasts for your factors, without having to manually change the column in the data frame. This approach means that the factors in the data table remain factors (allowing character values), avoiding the confusion that can result when factors are manually transformed into numeric predictors that are stored in the data table.

10	1	0	1	105	0	105
11	1	0	0	0	0	0
12	1	1	0	55	55	0
13	1	0	1	-5	0	-5
14	1	1	0	45	45	0
15	1	1	0	-10	-10	0

Note that the formula that we use to define our LM contains only all the *known* or *observable* variables, both the outcome variable and the input variables. The *unknown latent* variables—the β s and σ —are not mentioned in the formula. That’s because those are the variables we’re asking to be *estimated*, e.g., one β for each predictor. The estimates that we get from fitting the LM are what is reported in the output of the LM if you print or summarize it.

Working through an example

Next, we’ll go through an example. We’ll use the data from Tan & Jaeger (2024). *Learning to understand an unfamiliar talker: Testing models of adaptive speech perception*. The data come from an experiment on adaptive speech perception, and contain both responses and reaction times from a two-alternative forced-choice task (2AFC). From the abstract of Tan & Jaeger:

[...] a few minutes of exposure can significantly reduce the processing difficulty listeners experience during initial encounters with an unfamiliar accent. How such adaptation unfolds incrementally, however, remains largely unknown, leaving basic predictions by theories of adaptive speech perception untested. [...] We begin to address these knowledge gaps in a novel incremental exposure-test paradigm. We expose US English listeners to shifted phonetic distributions of word-initial stops (e.g., “dill” vs. “till”), while incrementally assessing cumulative changes in listeners’ perception. [...]

The following figure illustrates the design of the study, consisting of three between-participant exposure conditions, crossed with nine within-participant blocks, each eliciting responses along the phonetic VOT continuum (measured at 12 points). VOT refers to voice onset timing, and is the primary cue to the contrasts like *dip* and *tip* in US English. It is the time between the opening of a speaker’s lips (release the air burst) and the start of vocal fold vibrations. As little as 5ms difference in that time can make the difference between hearing the word *dip* (lower VOT) and the word *tip* (higher VOT). The three between-participant condition differed only in the distribution of VOT during exposure blocks.

Specifically, the green condition has the exact same distribution as the red condition, but shifted 10ms to the right; the blue condition is shifted +40ms to the right. This design makes it possible to assess





Figure 2: The same simple 2AFC task on all trials. Participants started each trial by clicking the green button. A recording played, and **participants had to answer which of the two words they heard.**



how participants' interpretation of the same physical VOT input during the gray test blocks changes as a function of the differently shifted VOT distributions during the exposure blocks. **This makes it possible to test whether listeners change their categorization functions over VOT as a function of recent input from a talker.**

121 participants were included for analysis, each providing 144 exposure trials (3 blocks) and 72 test trials (6 blocks). Participants were unaware of the exposure-test structure of the experiment, and the procedure for exposure and test trials was identical: participants saw two words on the screen (e.g., “dip” and “tip”). Participants then heard a recording of a single word and then had to click the word they heard. All recordings were spoken by the same voice. Here **we use only the data from the six gray test blocks**, leaving 26352 observations.

The data

Since we want to avoid violations of the assumption of independence, we use the data from a single participant from the Shift40 condition, living us with $6 * 12 = 72$ observations. Here's some of the data of that participant. If you downloaded this R markdown document, feel free to change the participant, and see how that changes the results.

	ParticipantID	Condition	Phase	Block	Trial	VOT	Response.isT	Response.RT
1	124	Shift40	test	1	1	50	TRUE	2325
2	124	Shift40	test	1	2	55	TRUE	1791
3	124	Shift40	test	1	3	25	FALSE	1445
4	124	Shift40	test	1	4	30	FALSE	1999
5	124	Shift40	test	1	5	15	FALSE	1453
6	124	Shift40	test	1	6	65	TRUE	1641
7	124	Shift40	test	1	7	40	TRUE	4900
8	124	Shift40	test	1	8	35	TRUE	1489
9	124	Shift40	test	1	9	-5	FALSE	3543
10	124	Shift40	test	1	10	5	FALSE	1587

Figure 3: Incremental exposure-test design. Histograms summarize distribution of voice onset time (VOT) in each block (dark bars: “d”-words; light: “t”-words). The three *between-participant* conditions (rows) differed only in the distribution of VOT during exposure blocks. VOTs in test blocks were identical within & across conditions.

11	124	Shift40	test	1	11	45	FALSE	5549
12	124	Shift40	test	1	12	70	TRUE	1211
13	124	Shift40	test	3	1	30	FALSE	1337
14	124	Shift40	test	3	2	45	FALSE	1788
15	124	Shift40	test	3	3	50	FALSE	3213

Example 1: Analyzing the effects of VOT on response times

We'll begin by analyzing the effects of VOT on response RTs. This makes response RTs our outcome variable, and VOTs the input variable. The first decision we need to make is how we'd like to code VOT. For example, we could treat VOT as a continuous predictor or as a factor with 12 distinct levels (the number of measurement locations along VOT). And, if we treat VOT as a continuous predictor, we could include only linear VOT or also quadratic VOT or higher order polynomials, provided we have enough data to avoid overfitting.⁷ Similarly, we could decide to include log-transformed VOT instead of VOT, or other transforms. As described above, the linearity assumption of LMs does not prevent us from applying non-linear transform to any of the input variables.

For the present example, we will treat VOT as a continuous predictor and only include its linear component.⁸

```
m <- lm(formula = Response.RT ~ 1 + VOT, data = dd)
summary(m)
```

Call:
lm(formula = Response.RT ~ 1 + VOT, data = dd)

Residuals:

Min	1Q	Median	3Q	Max
-749.9	-451.7	-287.1	2.4	3685.0

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2205.812	191.953	11.491	<2e-16 ***
VOT	-7.595	4.556	-1.667	0.1 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 856.6 on 70 degrees of freedom
Multiple R-squared: 0.03818, Adjusted R-squared: 0.02444
F-statistic: 2.779 on 1 and 70 DF, p-value: 0.1

This output can be split into four parts. At the top, the model call is repeated, which is convenient, for example, when we look at a stored model object that we created some time ago. Below it, we see a statistical summary of the residuals.⁹ This summary can be helpful in identifying violations of the assumption of normality. Here, for example, there is clear summary that the residuals are *not* normally distributed: the minimum/maximum and quantiles are not symmetrical around zero. This is to be expected—but often

⁷ As a rule of thumb, the number of coefficients in the LM should be less than $1/20$ th of the number of data points, $k < n/20$. But that's just a rule of thumb, and details depend on how balanced the data is with regard to the distribution of predictor values.

⁸ When we treat a predictor x as continuous, we making the assumption that a change in its value of 1 unit has the same effect on the outcome y , regardless of where along x that change is added. That follows from the linearity assumption mentioned above.

⁹ **Residuals** are the prediction error of the LM, the difference between the actual and the predicted outcome. Each row of the model matrix X corresponds to a prediction for that row, and thus has a corresponding residual ϵ .

ignored—for RTs, which are known to be right-skewed with heavy right tails (due to the boundedness of RTs).¹⁰

The third part of the output summarizes the coefficient estimates ($\hat{\beta}$), their standard errors, the resulting t -statistic, and its p -value. The t -statistic is a signed measure of the estimates distance from 0 in terms of its own standard error ($t = \frac{\hat{\beta}}{SE_{\hat{\beta}}}$).¹¹ The p -value is the result of a Student's t -test with the residual DFs of the LM as DFs for the test. E.g., for the effect of VOT, $t(70) = -1.6669391$. And if we look the probability density function for the Student t distribution to ask how probable this t -value is for $DF = 70$, we'll get 0.0999951, which—after rounding—matches the p -value reported in the LM for the effect of VOT.

Finally, the bottom part of the LM output provides multiple measures of the goodness of fit of the LM. Of these, the R^2 tells us how much of the variance in the outcome y is accounted for by the model (*for this sample*). Notably, an LM can have one or more statistically significant effects even if the model overall accounts for very little variance in our observed outcome. Just as the size of an effect is not to be confused with its theoretical relevance, the R^2 of a model is not to be confused with whether it is a good model. Many effects of high theoretical relevance might be very small and difficult to detect.¹² The F -statistic (here $F(1, 70) = 2.779$) tests whether the predictors included in the model *together* provide for a statistically significantly better fit, compared to an LM with just the intercept (the “null model”). The DFs for this test will always be the number of predictors in the model. We will return to this F -test when we talk about model comparison.

Write-up

For this particular example, there isn't much to report. Note that we wouldn't usually care about the intercept, and rarely report it. So we might write:

To test our hypothesis, we used linear regression to regress the participant's reaction time against VOT. The effect of VOT was not significant ($\hat{\beta} = -7.6$, $t = -1.67$, $p \geq .1$).

A few considerations:

- **Use $\hat{\beta}$ when we describe coefficient estimates and write up results.** This emphasizes that these are *estimates* of the assumed population parameters (β).
- Try to resist language that suggests significance even when it was not found, as in “... did not reach significance”.

¹⁰ For review and large-scale statistical simulations that assess the consequences of using LMs and its repeated-measures extension LMMs for the analysis of RTs, see Burchill & Jaeger (2024)-JML.

¹¹ This is just the standard definition of a t -value as a measure of how different an estimate is from some target value. Here, the target value is 0 since we want to know whether the coefficient for a predictor is different from 0. If it is, we say that the predictor “has an effect” (on the outcome).

¹² Of course, if our goal is to predict a certain human behavior—like an election, or whether someone will choose the red or blue pill—we might indeed only be interested in models that account for much of the observed variance. But keep in mind that we routinely aim to massively *reduce* natural variance in our experiments, e.g., by removing other influences. So be careful not to equate variance explained in a particular experiment with variance in human behavior explained.

- Do not take p -values to indicate effect strength. It is ok to write things like “highly significant effect” but that is not the same as a “large effect”. Similarly, a “large effect”—in whatever way that is determined—is not necessarily the same as a “theoretically particularly important” effect.

Questions to ponder in class

- How would you interpret the above output?
- How would we expect the output to change if we centered VOT by subtracting the mean VOT from each VOT value before entering it into the model? Which estimates of the model would change? Which ones would stay the same?
- How about if we also standardized VOT (by dividing it by its standard deviation)?
- How could we test whether the slope of VOT changes across Blocks?

Bonus question:

- How could we test a somewhat more interesting hypothesis that more ambiguous VOTs—VOTs that are more likely to elicit an equal proportion of “d”- and “t”-responses—are responded to more slowly?

Fitting

Consider the following fitted LM from the previous section:

```
Call:
lm(formula = Response.RT ~ 1 + VOT, data = dd)

Residuals:
    Min       1Q   Median       3Q      Max
-749.9 -451.7 -287.1    2.4 3685.0

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 2205.812   191.953   11.491  <2e-16 ***
VOT          -7.595     4.556   -1.667    0.1 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 856.6 on 70 degrees of freedom
Multiple R-squared:  0.03818,    Adjusted R-squared:  0.02444
F-statistic: 2.779 on 1 and 70 DF,  p-value: 0.1
```

How did the model fitting algorithm arrive at the estimates and what do they present? What do the standard errors (SEs) of the estimates represent? Intuitively, we know that the estimates are the

“best-fitting” estimates for the coefficients. What this means depends on how the goodness of fit is calculated.

One way to find the best-fitting estimates for the coefficients of an LM is to minimize the total prediction error. For each observation in our data, the prediction error or *residual* of the model m is defined as the difference between the predicted value of y , $E_m[y]$, and the actually observed value of y , i.e., $residual = y - E_m[y]$. Since we don’t want negative and positive prediction errors (resulting from over- and under-predicting y , respectively) to cancel each other out, we cannot simply sum those prediction errors to get an estimate of how well the model fits our data. Instead, we sum the squared residuals, giving us the residual sum of squares (RSS).¹³ Finding the best-fitting estimates then becomes the identical to finding the estimates that minimize the RSS.

To better understand what happens when we find the best-fitting estimates by minimizing the RSS, let’s define a few function that help use deliver predictions just like the final LM presented above does, but for arbitrary intercept values a and arbitrary slope values b for the effect of VOT. We also define a function that gets the prediction error for each observation in our data, and a function that sums up the squared prediction error:

```
get_prediction <- function(a, b, data = dd) {
  a + b * dd$VOT
}

get_prediction_error <- function(a, b, data = dd) {
  dd$Response.RT - get_prediction(a, b, data = data)
}

get_RSS <- function(a, b, data = dd) {
  sum(get_prediction_error(a, b, data = data)^2)
}
```

Now we can use these functions to visualize the predictions and calculate the RSSs for different hypotheses about a and b . You can use Acrobat Viewer to animate through the following examples (PDF system preview typically won’t work):

We can also visualize the RSS—i.e., the function we’re trying to minimize—depending on the a wide range of values for a and b . This illustrates two things. First, the estimates provided by the LM at the start of this section indeed minimize the RSS of the model. Second, RSS are a convex function of the parameters in the model (here a and b), with a single optimum. This very appealing property of LMs allows for efficient fitting algorithms that are *guaranteed* to converge within an arbitrary finite number of steps against the optimal parameters. Basically, it means that going downhill is bound to get you closer to the optimum, rather than stuck in a local minimum (as can be the case for more complex optimization problems, e.g., for many

¹³ One could instead use any other function that would result in only positive prediction errors (e.g., taking the *absolute* prediction error). However, using the squared residuals has both appealing properties, and it turns out to be identical—for LMs—to maximizing the likelihood of the data.

Figure 4: Raw RTs for one participant across the six test blocks (points), along with prediction line for different intercepts a and slopes b for the model $RT \sim 1 + VOT$. Compare how the RSS changes, depending on a and b .

types of non-linear models).




Figure 5: RSS for the model $RT \sim 1 + VOT$, depending on its intercept a and slope b . The white point shows that location of the estimates given in the LM output at the beginning of this section.

Here's an example of how one can use a general optimizer to find the best-fitting estimates for an LM. The example uses R's `optim()` function but you can try to implement it in your preferred programming language.¹⁴ You can compare its output to the estimates reported at the start of the section:

```
# let's store the optimizer's steps in a data.frame
d.trace <- tibble(.rows = 0)

o <-
  optim(
    # Starting values for a and b
    par = c(0, 0),
    fn = function(x) {
      rss <- get_RSS(x[1], x[2]);
      # Uncomment this line to see how optim() searches the parameter space
      # message(paste0("a = ", x[1], ", b = ", x[2], "; RSS = ", rss, "\n"))
      # <<- assigns values to *global* variables (outside the function)
      d.trace <- d.trace %>% bind_rows(tibble(a = x[1], b = x[2], RSS = rss))
    }
  )
```

¹⁴ By default, `optim()` uses Quasi-Newtonian optimization. This works just fine for the present purpose.

```

    return(rss)
  })

# print optimal parameters found by optim()
print(o$par)

[1] 2205.921965    -7.594757

```

Apply what you learned

While minimization of RSS works well for LMs, there is a more generally applicable approach that yields the same results for LMs but also can be used for the fitting of many others types of models, *likelihood maximization*. Under this approach, we find the best-fitting estimates that maximize the likelihood of the observed outcomes y . Recall that the LM assumes y to be normally distributed around the predicted $E[y]$. This means we can calculate the *likelihood* of each observation under the hypothesized normal distribution. This normal distribution is a function of the hypothesized values for all the model's coefficients (a and b in the example above) *and* the value hypothesized for *sigma*. For any hypothesis about a , b , and σ , we can get the density of an observation y under $\mathcal{N}(a + bx, \sigma)$. By multiplying the densities of all observations, we can get the data likelihood under the hypothesized values for a , b , and σ . In practice, we instead sum the log-transformed likelihoods. This is numerically more stable, and the x that maximizes $\log(f(x))$ is the same that maximizes $f(x)$.

- Define a function that calculates the sum of log-densities for the model $RT \sim 1 + VOT$ given a set of arguments a , b , and σ .¹⁵
- Make a figure like the one shown above for RSS. For this figure set σ to the σ from the LM at the start of this section.¹⁶ We don't expect the likelihood values to be identical to the RSS values—they are on completely different scales. But do the surfaces have similar shapes (curvatures)?
- To further understand the relationship between OLS and ML fitting of linear regression with normally distributed residuals, see [this presentation, slides 99-122](https://faculty.washington.edu/cadolph/mle/topic2.p.pdf) (it's a lot of slides but that's just because each derivation step is presented as a separate slide).

Session info

```

devtools::session_info()

- Session info -----
setting  value

```




Figure 6: Same as above but zoomed in and showing the path of the optimizer superimposed. Labels show step of optimizer

¹⁵ In R, `dnorm()` will give you the density of an observation or vector of observations y under a Normal with the specified mean and standard deviations. Use the optional argument `log = TRUE` to get log densities.

¹⁶ In R, use `sigma()` to extract σ from an LM.

```

version R version 4.3.2 (2023-10-31)
os      macOS Sonoma 14.5
system  aarch64, darwin20
ui      X11
language (EN)
collate en_US.UTF-8
ctype   en_US.UTF-8
tz      America/New_York
date    2024-09-06
pandoc  3.1.11 @ /Applications/RStudio.app/Contents/Resources/app/quarto/bin/tools/aarch64/ (via rmarkdown)

```

```

- Packages -----
package * version      date (UTC) lib source
backports 1.5.0      2024-05-23 [1] CRAN (R 4.3.3)
broom     * 1.0.6      2024-05-17 [1] CRAN (R 4.3.3)
cachem    1.1.0      2024-05-16 [1] CRAN (R 4.3.3)
cli       3.6.2      2023-12-11 [1] CRAN (R 4.3.1)
colorspace 2.1-0      2023-01-23 [1] CRAN (R 4.3.0)
crayon    1.5.2      2022-09-29 [1] CRAN (R 4.3.0)
devtools  2.4.5      2022-10-11 [1] CRAN (R 4.3.0)
digest    0.6.35     2024-03-11 [1] CRAN (R 4.3.1)
dplyr     * 1.1.4      2023-11-17 [1] CRAN (R 4.3.1)
ellipsis  0.3.2      2021-04-29 [1] CRAN (R 4.3.0)
evaluate  0.23       2023-11-01 [1] CRAN (R 4.3.1)
fansi     1.0.6      2023-12-08 [1] CRAN (R 4.3.1)
farver    2.1.2      2024-05-13 [1] CRAN (R 4.3.3)
fastmap   1.2.0      2024-05-15 [1] CRAN (R 4.3.3)
forcats   * 1.0.0      2023-01-29 [1] CRAN (R 4.3.0)
fs        1.6.4      2024-04-25 [1] CRAN (R 4.3.1)
generics  0.1.3      2022-07-05 [1] CRAN (R 4.3.0)
gganimate * 1.0.9      2024-02-27 [1] CRAN (R 4.3.1)
ggplot2   * 3.5.1      2024-04-23 [1] CRAN (R 4.3.1)
ggthemes  * 5.1.0      2024-02-10 [1] CRAN (R 4.3.1)
gifski    1.12.0-2   2023-08-12 [1] CRAN (R 4.3.0)
glue      1.7.0      2024-01-09 [1] CRAN (R 4.3.1)
gtable    0.3.5      2024-04-22 [1] CRAN (R 4.3.1)
hms       1.1.3      2023-03-21 [1] CRAN (R 4.3.0)
htmltools 0.5.8.1    2024-04-04 [1] CRAN (R 4.3.1)
htmlwidgets 1.6.4      2023-12-06 [1] CRAN (R 4.3.1)
httpuv    1.6.15     2024-03-26 [1] CRAN (R 4.3.1)
knitr     * 1.47       2024-05-29 [1] CRAN (R 4.3.3)
labeling  0.4.3      2023-08-29 [1] CRAN (R 4.3.0)
later     1.3.2      2023-12-06 [1] CRAN (R 4.3.1)
lifecycle 1.0.4      2023-11-07 [1] CRAN (R 4.3.1)
lubridate * 1.9.3      2023-09-27 [1] CRAN (R 4.3.1)
magrittr  * 2.0.3      2022-03-30 [1] CRAN (R 4.3.0)
MASS      * 7.3-60.0.1 2024-01-13 [1] CRAN (R 4.3.1)
memoise   2.0.1      2021-11-26 [1] CRAN (R 4.3.0)
mime      0.12       2021-09-28 [1] CRAN (R 4.3.0)
miniUI    0.1.1.1    2018-05-18 [1] CRAN (R 4.3.0)
munSELL    0.5.1      2024-04-01 [1] CRAN (R 4.3.1)
pillar    1.9.0      2023-03-22 [1] CRAN (R 4.3.0)
pkgbuild  1.4.4      2024-03-17 [1] CRAN (R 4.3.1)
pkgconfig 2.0.3      2019-09-22 [1] CRAN (R 4.3.0)
pkgload   1.3.4      2024-01-16 [1] CRAN (R 4.3.1)
prettyunits 1.2.0      2023-09-24 [1] CRAN (R 4.3.1)
profvis   0.3.8      2023-05-02 [1] CRAN (R 4.3.0)
progress  1.2.3      2023-12-06 [1] CRAN (R 4.3.1)
progressr 0.14.0     2023-08-10 [1] CRAN (R 4.3.0)
promises  1.3.0      2024-04-05 [1] CRAN (R 4.3.1)
purrr     * 1.0.2      2023-08-10 [1] CRAN (R 4.3.0)

```

R6	2.5.1	2021-08-19	[1]	CRAN (R 4.3.0)
Rcpp	1.0.12	2024-01-09	[1]	CRAN (R 4.3.1)
readr	* 2.1.5	2024-01-10	[1]	CRAN (R 4.3.1)
remotes	2.5.0	2024-03-17	[1]	CRAN (R 4.3.1)
rlang	1.1.4	2024-06-04	[1]	CRAN (R 4.3.3)
rmarkdown	2.27	2024-05-17	[1]	CRAN (R 4.3.3)
rstudioapi	0.16.0	2024-03-24	[1]	CRAN (R 4.3.1)
scales	1.3.0	2023-11-28	[1]	CRAN (R 4.3.1)
sessioninfo	1.2.2	2021-12-06	[1]	CRAN (R 4.3.0)
shiny	1.8.1.1	2024-04-02	[1]	CRAN (R 4.3.1)
stringi	1.8.4	2024-05-06	[1]	CRAN (R 4.3.1)
stringr	* 1.5.1	2023-11-14	[1]	CRAN (R 4.3.1)
tibble	* 3.2.1	2023-03-20	[1]	CRAN (R 4.3.0)
tidyr	* 1.3.1	2024-01-24	[1]	CRAN (R 4.3.1)
tidyselect	1.2.1	2024-03-11	[1]	CRAN (R 4.3.1)
tidyverse	* 2.0.0	2023-02-22	[1]	CRAN (R 4.3.0)
timechange	0.3.0	2024-01-18	[1]	CRAN (R 4.3.1)
tinytex	0.51	2024-05-06	[1]	CRAN (R 4.3.1)
tufte	* 0.13	2023-06-22	[1]	CRAN (R 4.3.0)
tweenr	2.0.3	2024-02-26	[1]	CRAN (R 4.3.1)
tzdb	0.4.0	2023-05-12	[1]	CRAN (R 4.3.0)
urlchecker	1.0.1	2021-11-30	[1]	CRAN (R 4.3.0)
usethis	2.2.3	2024-02-19	[1]	CRAN (R 4.3.1)
utf8	1.2.4	2023-10-22	[1]	CRAN (R 4.3.1)
vctrs	0.6.5	2023-12-01	[1]	CRAN (R 4.3.1)
viridisLite	0.4.2	2023-05-02	[1]	CRAN (R 4.3.0)
withr	3.0.0	2024-01-16	[1]	CRAN (R 4.3.1)
xfun	0.44	2024-05-15	[1]	CRAN (R 4.3.3)
xtable	1.8-4	2019-04-21	[1]	CRAN (R 4.3.0)
yaml	2.3.8	2023-12-11	[1]	CRAN (R 4.3.1)

[1] /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library
